

# Attention Models and Memory Networks

Injung Kim  
Handong Global University

# Agenda

---




- Introduction
- Encoder–Attention–Decoder Models
- Memory Networks
- Transformer
- Visual Attention Models
- Q&A

# Deep Learning Architectures

- Deep Neural Networks (DNN)

- MLP, RBM, DBN, DBM, ...




- Convolutional neural networks (CNN)

- Image processing
- Learns position invariant local features


- Recurrent neural networks (RNN)

- Sequence processing (text, speech, etc.)
- Recurrent connection (short-term memory)




- Deep generative models

- GAN, VAE, autoreg. models, flow models, etc.



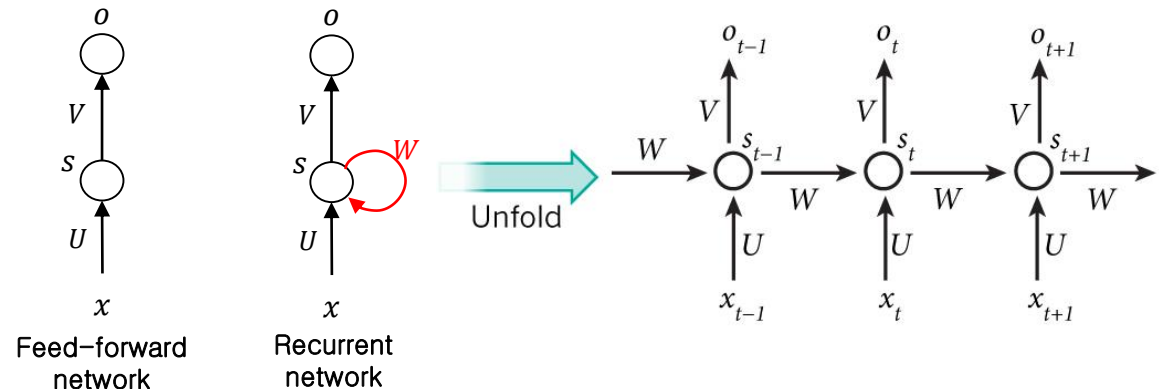

- Attention models

- Encoder-attention-decoder, transformer, visual attention models
- NLP, ASR/TTS, image processing
- Closely related to [memory network](#)




# Recurrent Neural Networks

- Recurrent neural network is a neural network specialized for processing **a sequence of values**  $x^{(1)}, x^{(2)}, \dots, x^{(\tau)}$ .
  - Neural networks with recurrent connection
  - **State of nodes** affect the output and the next state
  - Model for dynamic process
  - Temporarily shared connections





# Encoder-Decoder Models

- Architecture for sequence-to-sequence mapping
  - Encoder: input sequence → **context vector**
  - Decoder: **context vector** (+ previous output) → new output
- The two modules are trained jointly to maximize the average of  $\log P(y^1, y^2, \dots, y^m | x^1, x^2, \dots, x^n)$



# Encoder-Attention-Decoder Models

- Decoder accesses context composed of weighted average of input states



# Attention Models [Bahdanau16]

- Output

$$p(y_i | y_1, \dots, y_{i-1}, \mathbf{x}) = g(y_{i-1}, s_i, c_i)$$

- Hidden state

$$s_i = f(s_{i-1}, y_{i-1}, c_i)$$

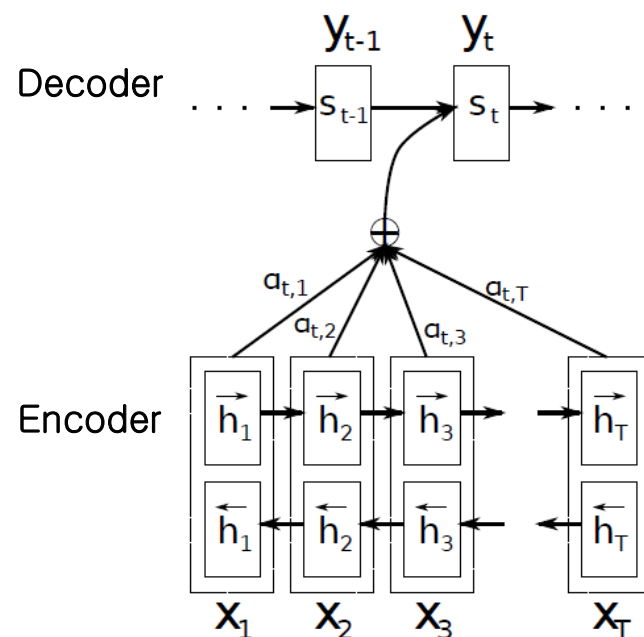
- Context vector

$$c_i = \sum_{j=1}^{T_x} \alpha_{ij} h_j$$

value

- Attention

$$\alpha_{ij} = \frac{\exp(e_{ij})}{\sum_{k=1}^{T_x} \exp(e_{ik})} \quad e_{ij} = a(\underbrace{s_{i-1}}_{\text{query}}, \underbrace{h_j}_{\text{key}})$$



# Agenda

---



- Introduction
- Encoder–Attention–Decoder Models
- Memory Networks
- Transformer
- Visual Attention Models
- Q&A



# Memory Networks



## ■ Motivation

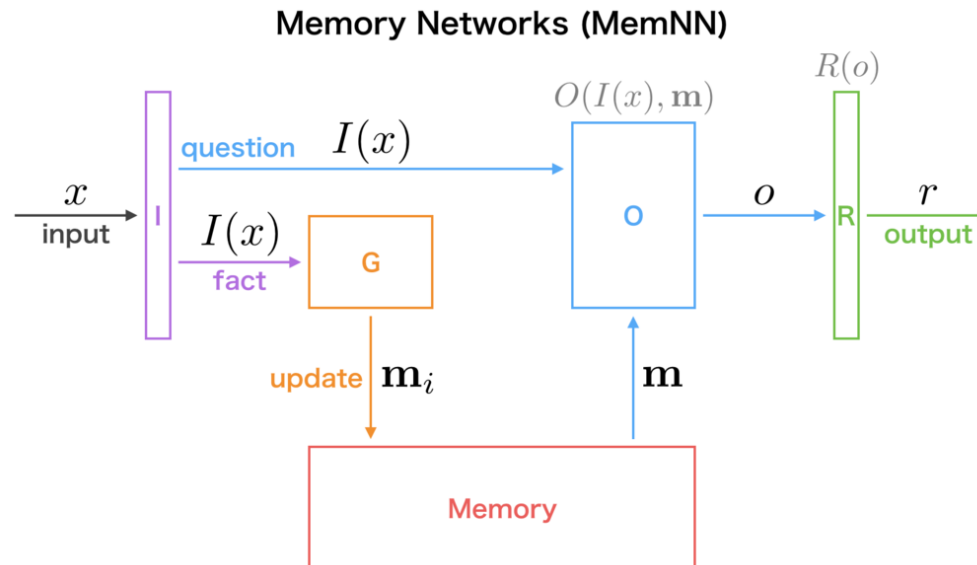
- Conventional encoder–decoder models only maps the input sequence to an output sequence
  - Question → Answer
- Some tasks require explicit memory
  - Read a long story. Then, answer questions about that.
- Conventional machine learning models lack **an easy way to read and write** to part of a (potentially very large) long-term memory component.

## ■ MemNN

- Neural network **with explicit memory**

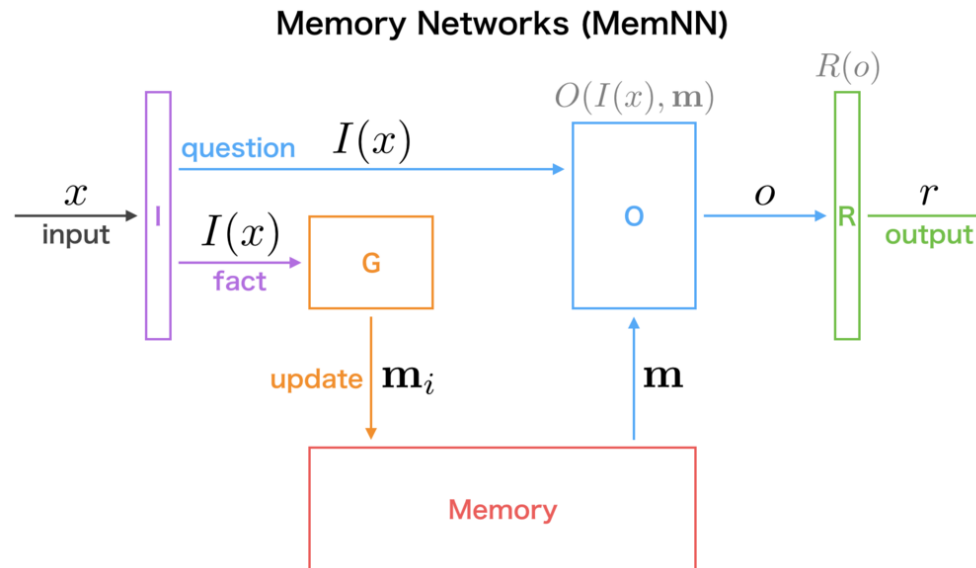
# Memory Networks

- Weston, et al., “Memory Networks,” Oct. 2014
  - Combine the successful learning strategies for **inference with a memory component** that can be read and written to.
  - Memory  $m$ : an array of objects indexed by  $m_i$
  - Four components (I, G, O and R)




# Memory Networks

- Four (potentially learned) components
  - I (input feature map): input  $\rightarrow$  feature
  - G (generalization): updates old memories given input
  - O (output feature map): (input + current memory state)  
 $\rightarrow$  (multiple) memory output
  - R (response): memory output  $\rightarrow$  response



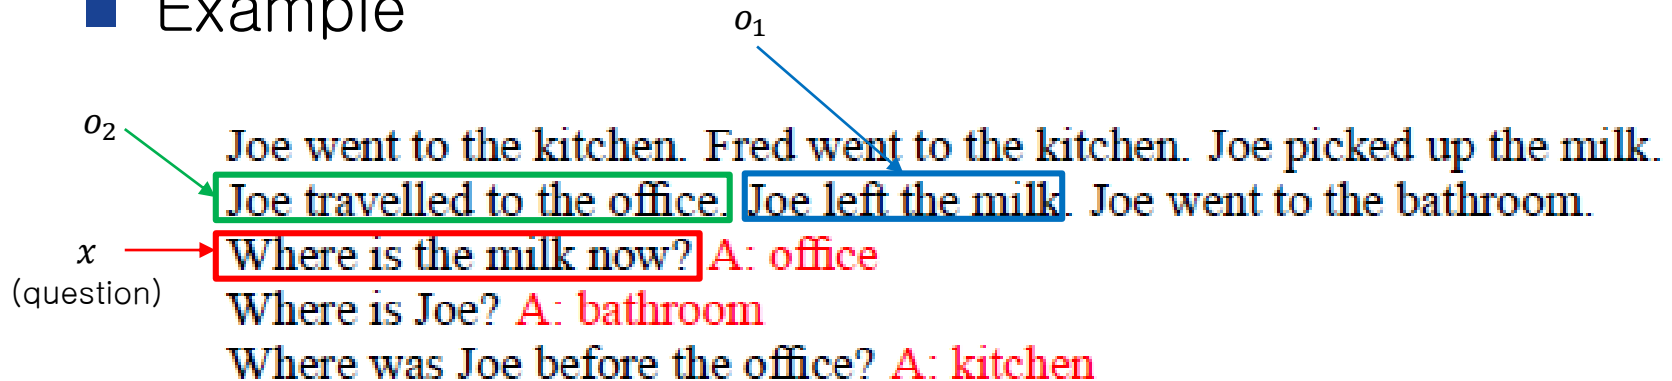
# Memory Networks

- Operation given an input  $x$ ,
  - The same for training/test time. But, model parameters are updated only in training time.
    1. Convert  $x$  to an internal feature representation  $I(x)$ .
    2. Update memories  $\mathbf{m}_i$  given the new input:  $\mathbf{m}_i = G(\mathbf{m}_i, I(x), \mathbf{m}), \forall i$ .
    3. Compute output features  $o$  given the new input and the memory:  $o = O(I(x), \mathbf{m})$ .
    4. Finally, decode output features  $o$  to give the final response:  $r = R(o)$ .

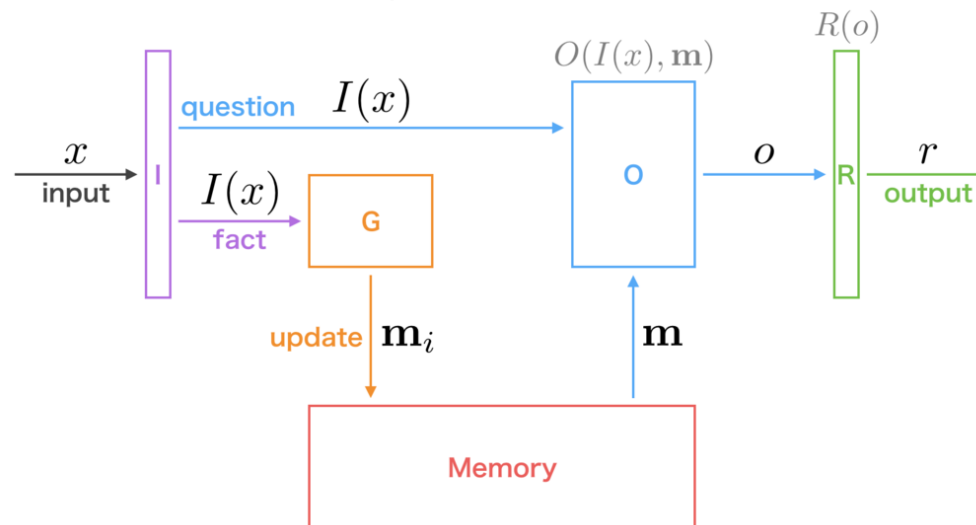


# Memory Networks

## ■ Example



Memory Networks (MemNN)



# Hard Addressing vs. Soft Addressing

## ■ Memory

- $M = (M_1, M_2, \dots, M_m)$ 
  - $M_i$  is memory content ( $n$ -dim vector)

## ■ Hard addressing ( $a$ : address)

- $r = M_a$
- $r = \sum_i w_i M_i$ 
  - $w$  is a one hot vector. ( $w_i = 1$  if  $i = a$ , and  $w_i = 0$ , otherwise).

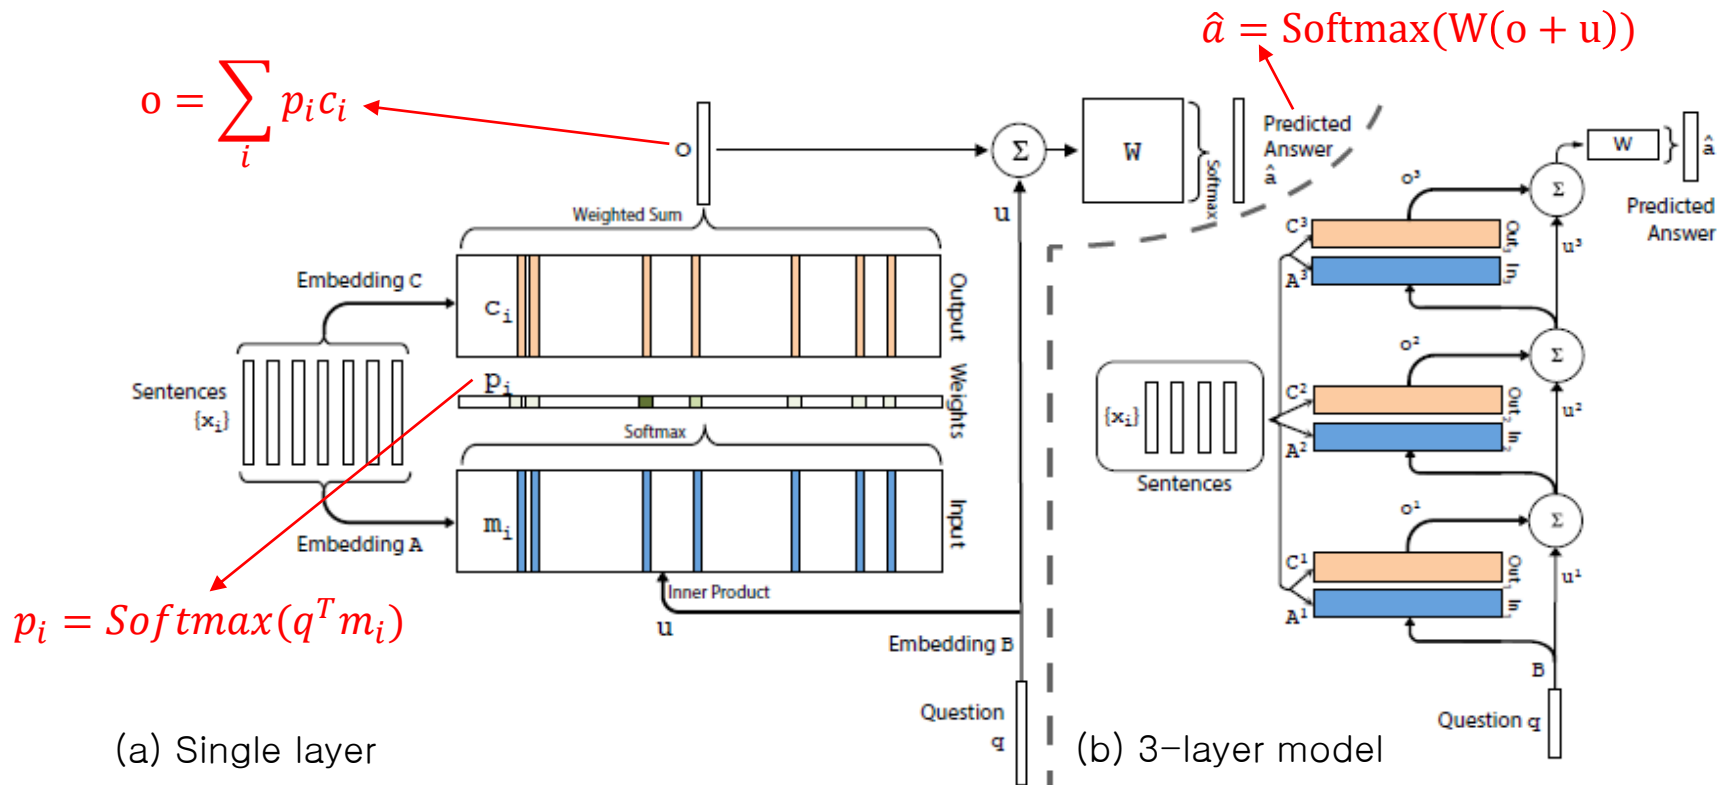
## ■ Soft addressing ( $w$ : soft address vector)

- $r = \sum_i w_i M_i$ 
  - $w$  is an arbitrary vector (usually,  $\sum_i w_i = 1$ )

$w_1$	$w_2$	$\dots$	$w_a$	$\dots$	$w_m$
$M_1$	$M_2$	$\dots$	$M_a$	$\dots$	$M_m$

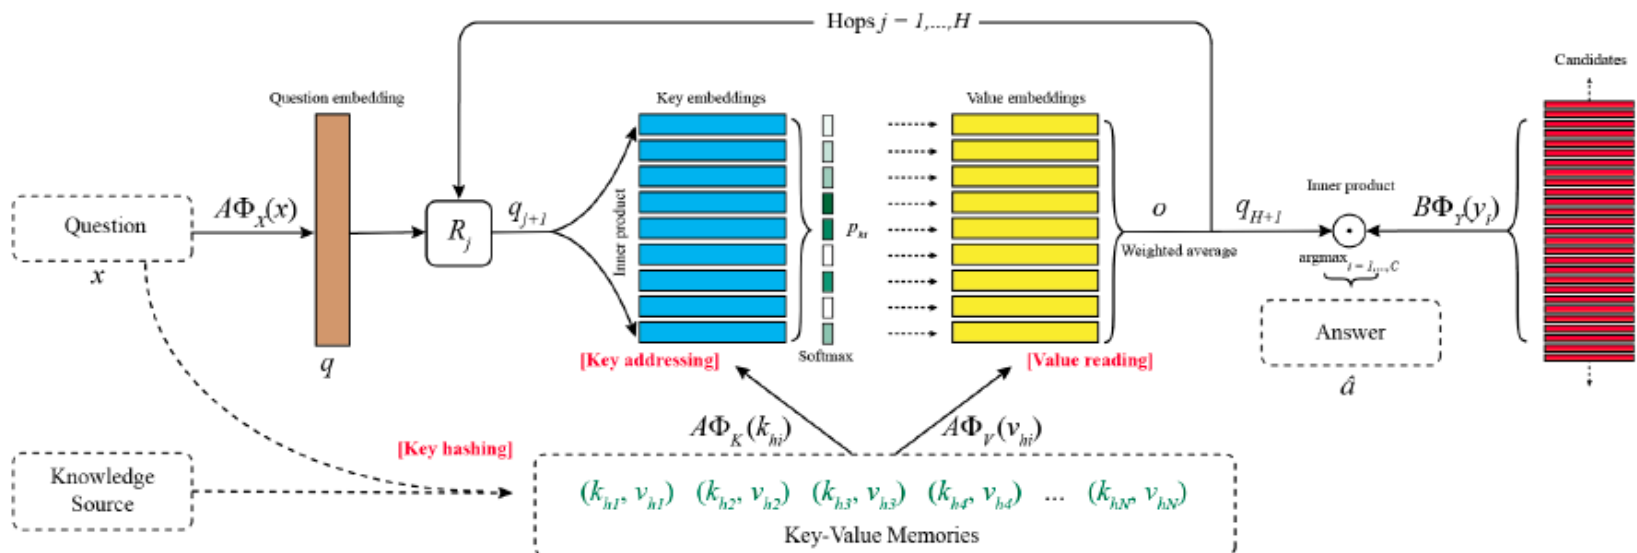
# End-to-End Memory Networks [Sukhbaatar15]

- Inputs: a set of inputs  $(x_1, x_2, \dots, x_n) + \text{query } (q)$
- Output: answer  $(a)$ 
  - Produced by multiple *hops*.



# Key-Value Memory Network

- Miller, et al., “Key-Value Memory Networks for Directly Reading Documents”, June, 2016.
  - KV-MemNN for question answering
  - Stores facts in key-value structured memory
    - Key: address relevant memories w.r.t. question
    - Value: the information the memory provides





# Content-based vs. Location-based Addressing

- Key-value memory
  - $M = (M_1, M_2, \dots, M_m), M_i = (K_i, V_i)$
- Accessing memory  $M$  with a weight vector  $w$ 
  - $r = \sum_i w_i V_i = wV$
- Content-based addressing
  - Determine  $w$  by similarity between query ( $Q$ ) and key ( $K_i$ )
    - $w_i^t = f(Q, K_i)$  (e.g.,  $\text{Softmax}(Q \cdot K_i)$ )
    - $r = \sum_i w_i V_i$
- Location-based addressing
  - Choose new address near the previous address
    - $w_i^t = w_i^{t-1} + \Delta w_i^t$
    - $\Delta w_i^t = g(Q, K_i, w_i^{t-1})$

# Attention Models [Bahdanau16]

- Output

$$p(y_i | y_1, \dots, y_{i-1}, \mathbf{x}) = g(y_{i-1}, s_i, c_i)$$

- Hidden state

$$s_i = f(s_{i-1}, y_{i-1}, c_i)$$

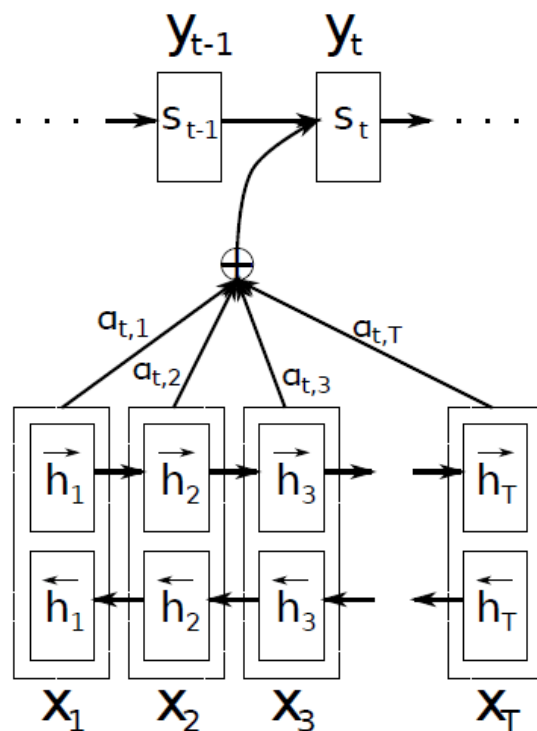
- Context vector

$$c_i = \sum_{j=1}^{T_x} \alpha_{ij} h_j$$

value

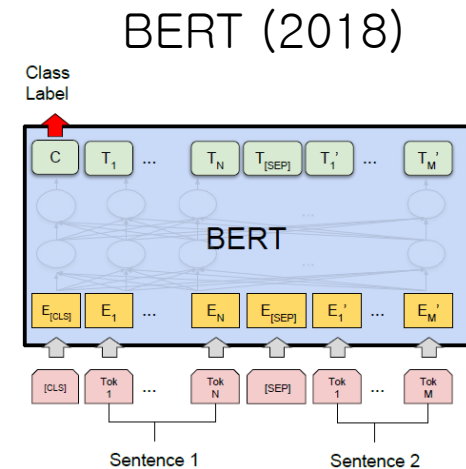
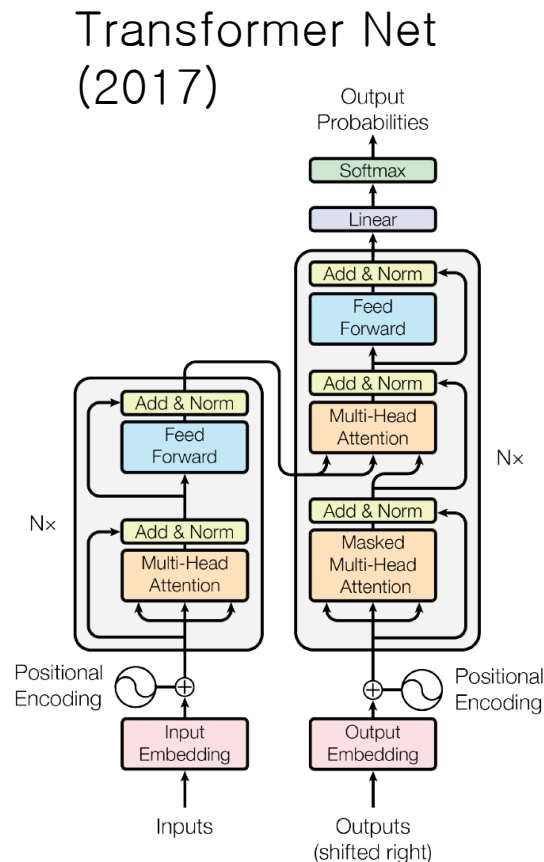
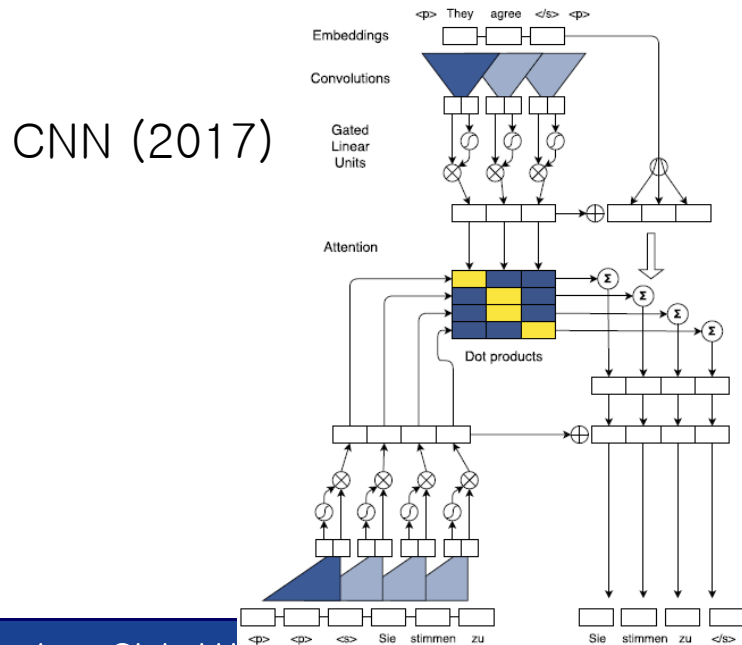

- Attention

$$\alpha_{ij} = \frac{\exp(e_{ij})}{\sum_{k=1}^{T_x} \exp(e_{ik})} \quad e_{ij} = a(\underbrace{s_{i-1}}_{\text{query}}, \underbrace{h_j}_{\text{key}})$$



# Natural Language Processing

- Encoder–Attention–Decoder models
  - RNN → CNN → Transformer → BERT, GPT1/2/3



# Agenda

---



- Introduction
- Encoder–Attention–Decoder Models
- Memory Networks
- Transformer
- Visual Attention Models
- Q&A

# Transformer

- A. Vaswan, “Attention Is All You Need” Jun. 2017
  - Attention mechanism instead of convolutional or recurrent units
  - Global dependencies between input and output
  - More parallelization
  - Reduced effective resolution
  - ➔ Counteract with multi-head attention

- References

- Illustrated-transformer (<http://jalamar.github.io/illustrated-transformer/>)
- The annotated transformer (<http://nlp.seas.harvard.edu/2018/04/03/attention.html>)

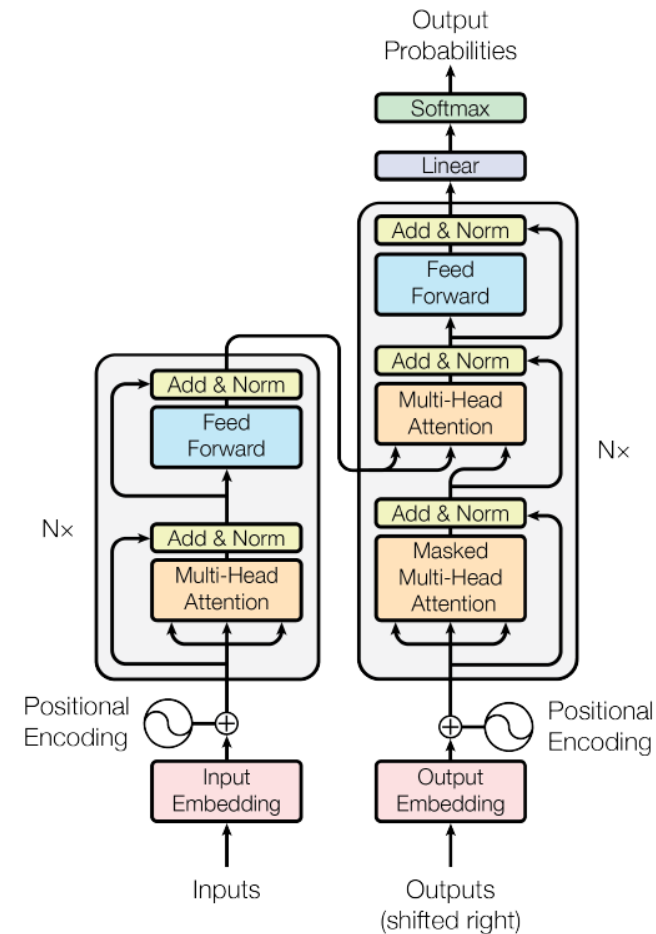
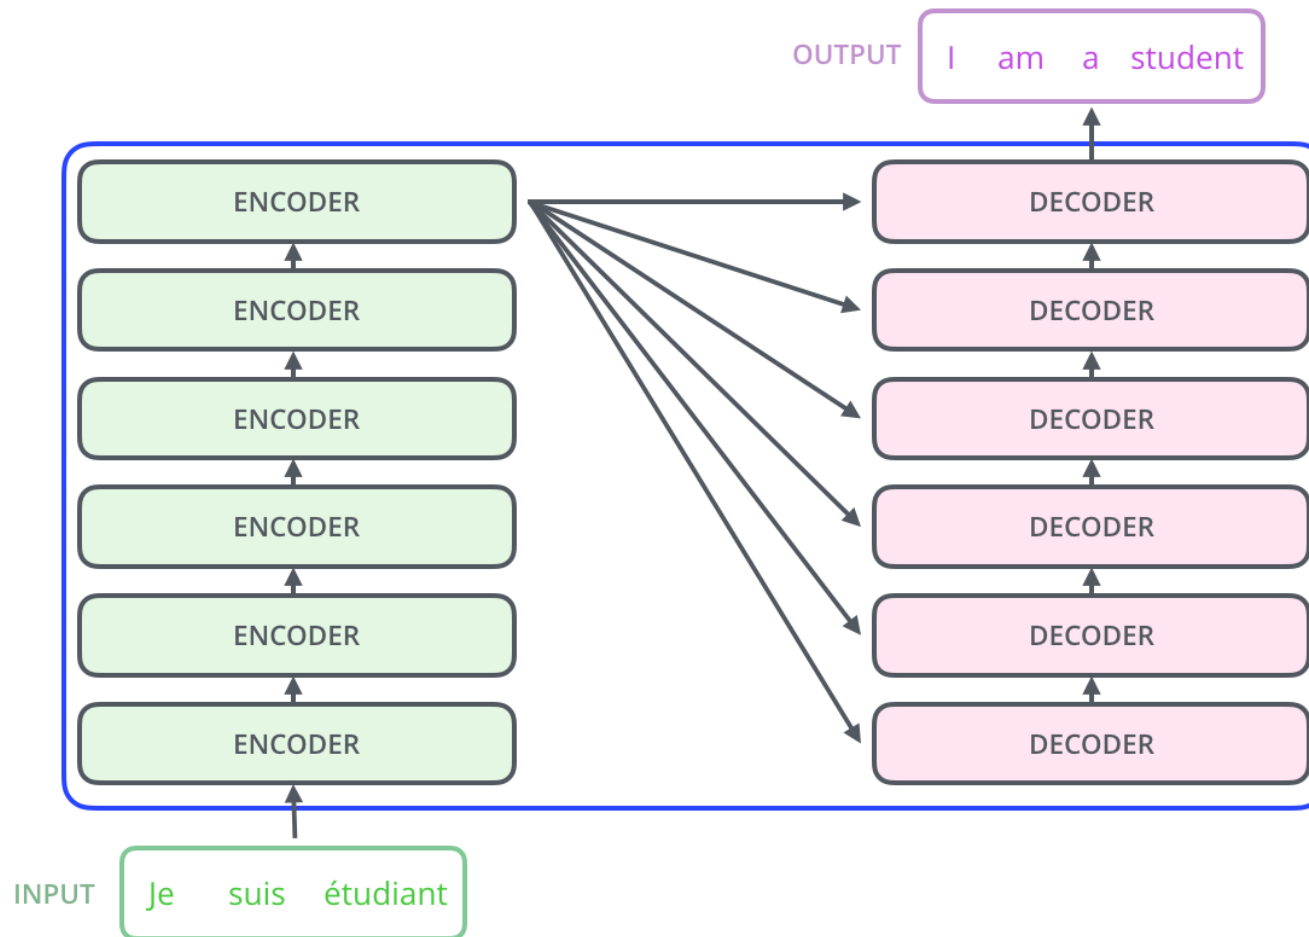


Figure 1: The Transformer - model architecture.

# Transformer

- Illustration of Transformer operation



# Transformer Encoder

- Encoder consists of  $N=6$  identical layers
  - 1<sup>st</sup> sublayer: multi-head self-attention
  - 2<sup>nd</sup> sublayer: position-wise FC
  - Residual connection + layer norm
$$\text{LayerNorm}(x + \text{sublayer}(x))$$
  - All sublayers  $d_{\text{model}} = 512$




Figure 1: The Transformer - model architecture.

# Transformer Decoder

- Decoder consists of  $N=6$  identical layers
  - Similar to encoder layers
  - Additional sublayer: performs multi-head attention over the output of the encoder stack.
  - Masked self-attention only to preceding positions

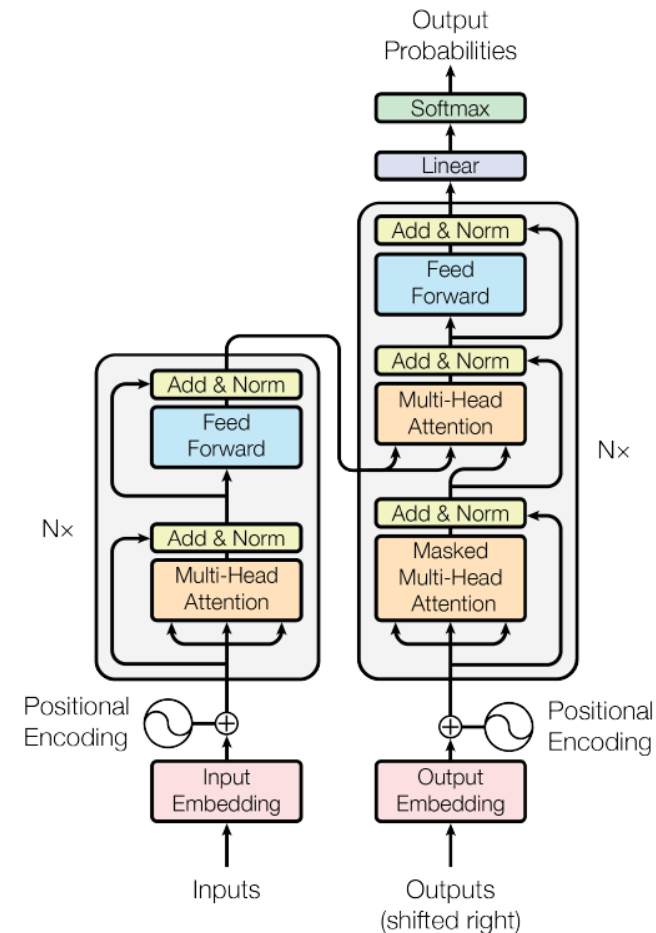


Figure 1: The Transformer - model architecture.



# Attention

- Scaled dot-product attention
  - Input: a query, a set of key-value pairs
  - Output: weighted sum of values

$$\text{Attention}(Q, K, V) = \text{softmax}\left(\frac{QK^T}{\sqrt{d_k}}\right) \cdot V$$


- Multi-header attention
  - Jointly attend to information from **different representation subspaces** at **different positions**

$$\text{MultiHead}(Q, K, V) = \text{Concat}(\text{head}_1, \dots, \text{head}_h)W^O$$


$$\text{head}_i = \text{Attention}(QW_i^Q, KW_i^K, VW_i^V)$$

$$\square h = 8, d_k = d_v = \frac{d_{\text{model}}}{h} = 64$$

Scaled Dot-Product Attention



Multi-Head Attention



# Agenda


---




- Introduction
- Encoder–Attention–Decoder Models
- Memory Networks
- Transformer
- Visual Attention Models
- Q&A

# Convolutional Neural Networks

- **Convolutional Neural networks (CNN)**: a class of deep feed-forward network designed to mimic human/animal visual systems
  - LeNet 1998, MNIST, CPU



- AlexNet 2012, ImageNet, GTX 580 x 2



# ResNet

- He, et.al, “Deep Residual Learning for Image Recognition”, Dec. 2015
  - Target function  $H(x) = F(x) + x$
  - Residual Function  $F(x) = H(x) - x$
  - Deep residual network contains 152 layers






Figure 2. Residual learning: a building block.



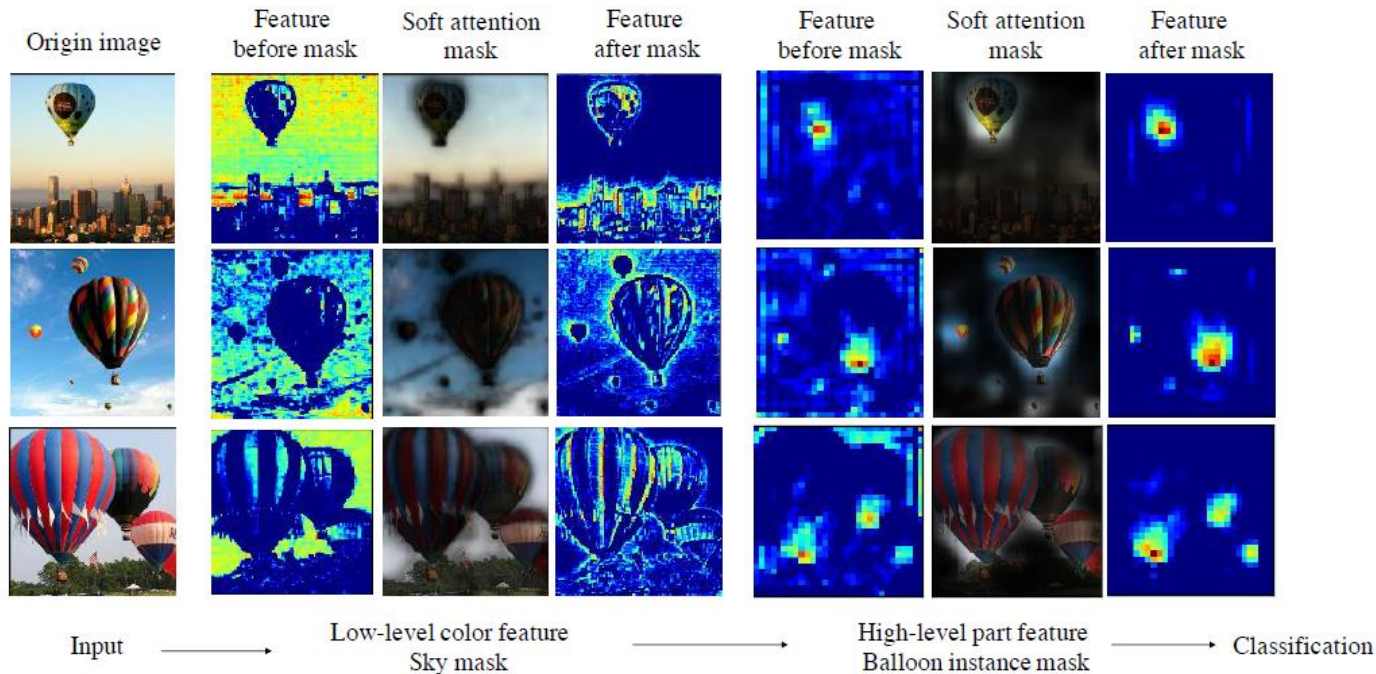
Bottleneck building block

# Residual Attention Network

- Wang, et al., “Residual Attention Network for Image Classification,” Apr. 2017.
  - Residual Attention Network built by stacking Attention Modules
  - Attention residual learning
  - Bottom-up top-down feedforward network (hourglass model)




Attention mechanism



# Residual Attention Network


- Wang, et al., “Residual Attention Network for Image Classification,” Apr. 2017.
  - Residual Attention Network built by stacking Attention Modules
  - Attention residual learning
  - Bottom-up top-down feedforward network (hourglass model)



- Resnet:  $H(x) = x + F(x)$
- RAN:  $T(x) + M(x)T(x) = (1 + M(x))T(x)$


# Residual Attention Network

- Wang, et al., “Residual Attention Network for Image Classification,” Apr. 2017.
  - Residual Attention Network built by stacking Attention Modules
  - Attention residual learning
  - Bottom-up top-down feedforward network (hourglass model)



# SENet


- Hu, et al., “Squeeze-and-Excitation Networks,” Sep. 2017.
  - Recalibrates channel-wise feature responses by explicitly modelling interdependencies between channels
  - “Attention along channels”





# Non-Local Nets

- Wang, et al., “Non-local Neural Networks,” Nov. 2017.
  - Learns relation among distant feature elements

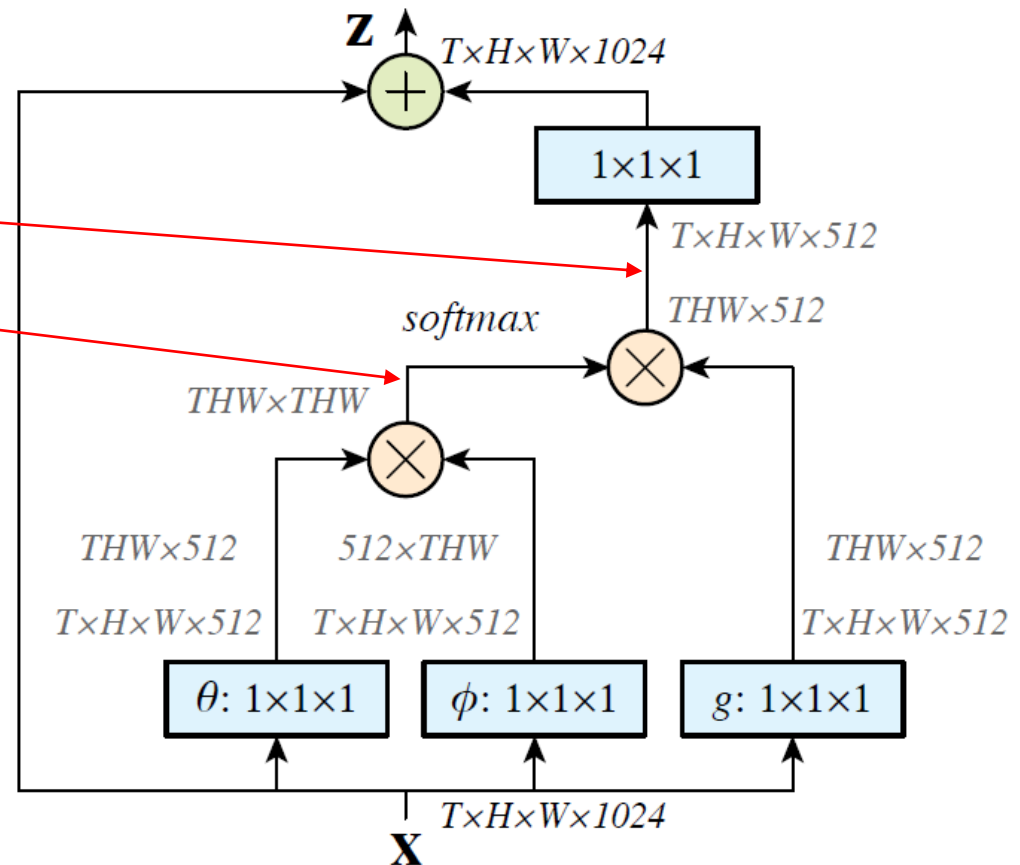


# Non-Local Nets

- Wang, et al., “Non-local Neural Networks,” Nov. 2017.

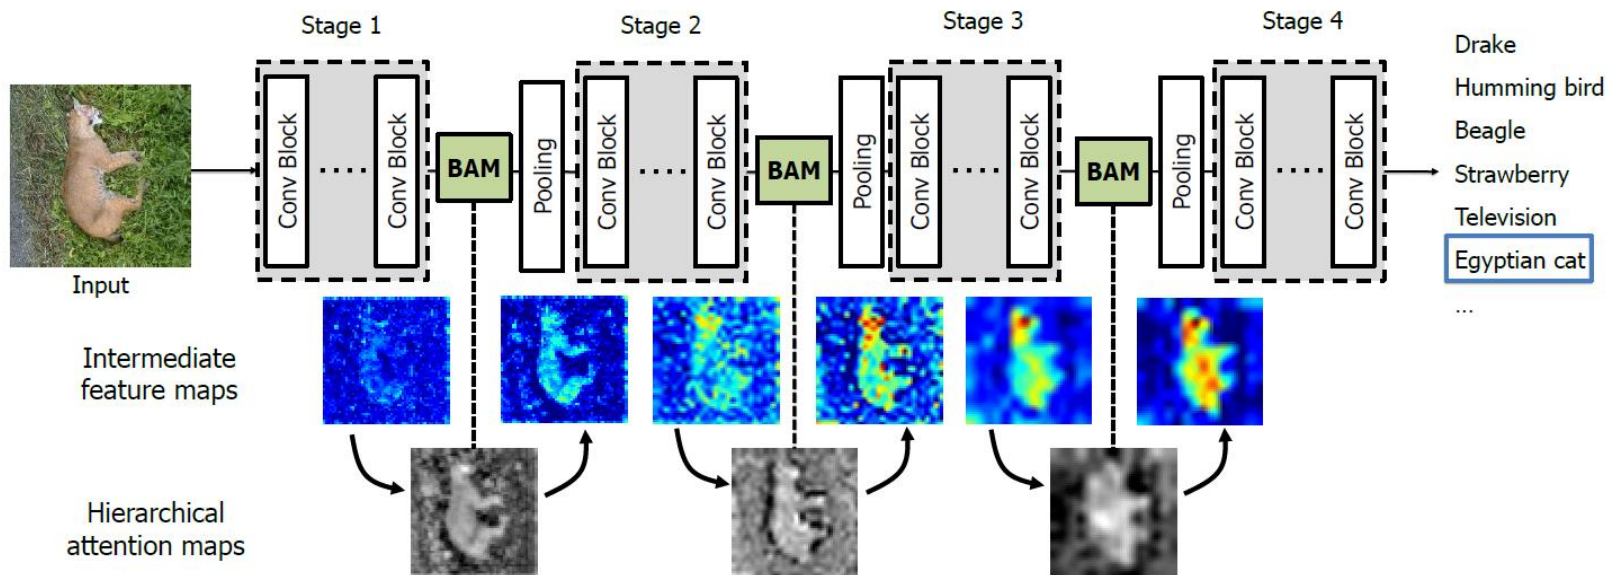
$$y_i = \frac{1}{C(\mathbf{x})} \sum_{\forall j} f(\mathbf{x}_i, \mathbf{x}_j) g(\mathbf{x}_j)$$

$$f(\mathbf{x}_i, \mathbf{x}_j) = e^{\theta(\mathbf{x}_i)^T \phi(\mathbf{x}_j)}$$

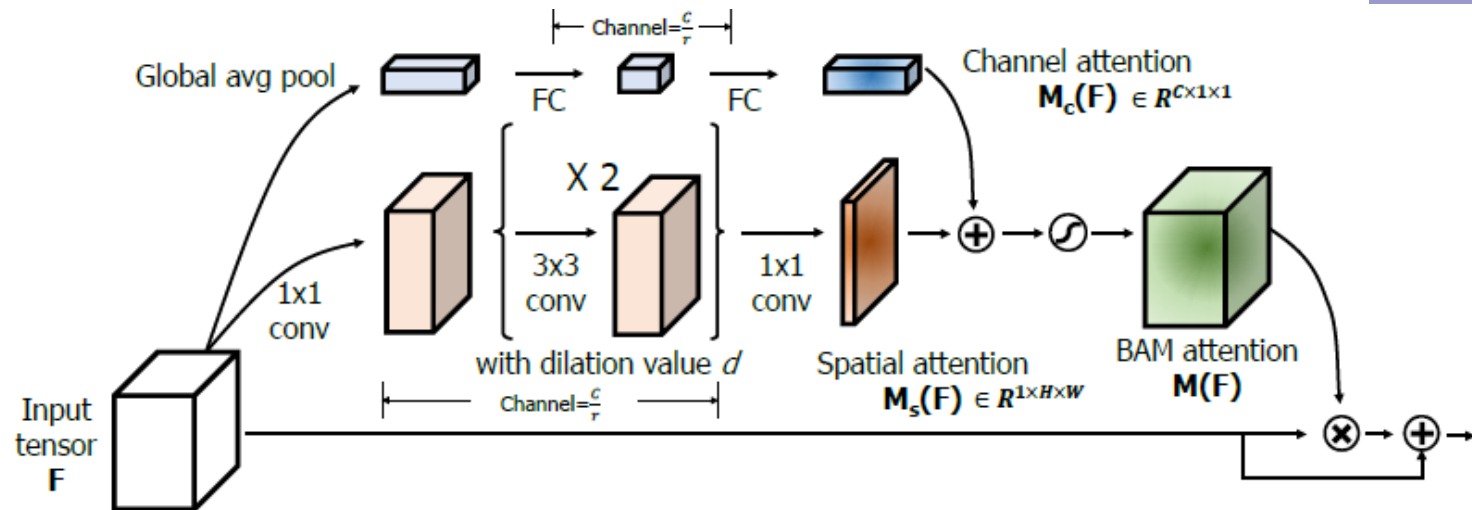


# BAM

- Park, et al., “BAM: Bottleneck Attention Module,” Jul. 2018.
  - Infers an attention map along two channel and spatial
  - Place BAM at each bottleneck of models where the downsampling of feature maps occurs.



# BAM



- Output feature map  $F' = F + F \oplus M(F)$
- 3D attention map  $M(F) = \sigma(M_c(F) + M_s(F))$
- Channel attention

$$\begin{aligned} \mathbf{M}_c(\mathbf{F}) &= BN(MLP(AvgPool(\mathbf{F}))) \\ &= BN(\mathbf{W}_1(\mathbf{W}_0 AvgPool(\mathbf{F}) + \mathbf{b}_0) + \mathbf{b}_1) \end{aligned}$$

- Spatial attention

$$\mathbf{M}_s(\mathbf{F}) = BN(f_3^{1 \times 1}(f_2^{3 \times 3}(f_1^{3 \times 3}(f_0^{1 \times 1}(\mathbf{F}))))))$$



Thank you  
for your attention!

